

The True Lethality Engine

Predicting the real danger of D&D 5e combat from 34,907 recorded encounters

Daniele · Francesca · Stefano · Antonietta

Statistical Machine Learning — Final Project, Sapienza Università di Roma, A.Y. 2025/2026

<https://dnd-appraising.streamlit.app> <https://github.com/danpinocontrollino/project-dnd>

Abstract

Dungeons & Dragons 5e rates every monster with a *Challenge Rating* (CR), a designer-assigned difficulty score that has never been validated against actual play. We validate it — and replace it. From 153,829 turn snapshots of real online play (the FIREBALL corpus) we build 34,907 encounter-level observations across 1,462 campaigns, parse offensive statistics directly out of monster statblock text, and train a calibrated, monotonicity-constrained XGBoost model of $\eta(x) = P(\text{party wins} \mid x)$. Inverting η by binary search yields the *True Lethality Level*: the party level at which a balanced party of four reaches a target win rate. Validation is grouped by campaign — the true unit of leakage — giving an honest held-out ROC-AUC of 0.657 (95% bootstrap CI [0.638, 0.674]) and Brier score 0.139 [0.134, 0.145] at an 83% base rate. Because observational play is contaminated by “DM mercy” (fights the damage math deems hopeless were still recorded as wins 84.5% of the time), the engine adds three explicit domain guards; the survival guard’s constants are calibrated externally against 490,640 simulated fight-to-the-death battles. Along the way, a ridge logistic model *wins* on observational AUC (0.655 vs. 0.614) yet is rejected: its confounded coefficients make it unusable for the interventional questions the application asks. The result is a live encounter appraiser that measures the rulebook against the data — with the rulebook’s own math implemented honestly.

1 Introduction

In D&D 5e a party of 3–6 heroes (levels 1–20) fights monsters run by a Dungeon Master (DM). The official difficulty currency is the Challenge Rating: a CR- k monster should challenge four level- k heroes, and the *Dungeon Master’s Guide* (DMG, p. 82) grades encounters by comparing multiplier-adjusted monster XP against per-level thresholds [2]. CR is a design intention: it compresses offense, defense, traits and action economy into one printed number, and its known failure modes (high-CR solos folding to action economy, deadly low-CR swarms, trait synergies) had never been quantified against play. We ask a precisely defined question instead:

What is the calibrated probability $\eta(x)$ that a given party wins a given encounter x — and at which party

level L does a balanced party of four reach a chosen target win rate (default 65%)?

L , found by binary search on η over levels [1, 20] (18 iterations, quarter-level rounding), is the **True Lethality Level**. Boundary cases get explicit verdicts rather than fabricated precision: *trivial* ($\eta \geq \text{target}$ at level 1) and *beyond deadly* ($\eta < \text{target}$ at level 20, reported with its ceiling).

Contributions. (i) an encounter-level dataset distilled from real play, with offensive statistics parsed from statblock text; (ii) a calibrated, shape-constrained model with campaign-grouped validation and bootstrap uncertainty; (iii) a diagnosis of estimand mismatch in observational game data (“DM mercy”) and a three-guard architecture, one guard calibrated by external simulation; (iv) a course-toolbox benchmark (penalized logistic, random-feature kernel machine, MMD test, Gaussian process) with a decision-theoretic model-selection lesson; (v) a deployed application that compares the model against the DMG’s own procedure, implemented faithfully.

2 Data

2.1 From chat logs to encounters

FIREBALL [1] collects real games played on Discord through the Avrae bot. Our ETL (`parse_fireball.py`) processes 1,471 JSONL logs (153,829 turn snapshots; 134,475 valid turns), infers outcomes from terminal HP states, drops 85,778 unresolved fights and 9,215 encounters with no bestiary match, and emits **34,907 resolved encounters** across **1,462 campaigns**. The party win base rate is **83.0%** — a number that shapes everything downstream. Each row has 29 raw columns: party level/size and four role flags (healer, tank, arcane, martial DPS, mapped from character classes), roster aggregates (count-weighted means for continuous statistics, maxima for binary traits and apex threats, pooled sums for the damage race), ten monster trait flags, XP totals, and the outcome. Party levels are clipped to [1, 20] and rosters to 30 (mass-summon logs reach 241 actors).

2.2 The Attack Potency problem

Monster manuals list HP and AC but no machine-readable offense. We parse it from the SRD stat-

block text itself: attack bonuses (+9 to hit), printed average damage (12 (2d6+5)), multiattack routines (“makes *three* attacks”), save DCs, and spell lists priced by a 30-spell PHB damage table. Per-action segmentation separates *sustained* DPR (weapon routine) from *burst* (the single scariest action): the Lich yields DPR 45 and burst 100 (*Power Word Kill*). Coverage: 321/762 monsters from real statblocks, the remainder imputed from the DMG’s design table by CR; parsed DPR is clamped to $[0.25\times, 3\times]$ of the design band. Regression tests caught two real bugs: versatile weapons double-counted (24 monsters) and dragon breath leaking into multiattack math (Adult Red Dragon $72.6 \rightarrow 58$; hand-computed truth $\approx 45\text{--}48$). Spot checks: Goblin 5 (true 5), Tarrasque 150 (true ≈ 148).

2.3 Playing fair with the rulebook

If we claim to beat the DMG we must implement it correctly. The application’s “by-the-book” verdict runs the full p. 82 procedure — total XP, the count multiplier ladder with the p. 83 party-size adjustment, and an inverse XP lookup to an equivalent single-monster CR — so that $6\times\text{CR}-1$ monsters read as one CR-6 encounter, not “CR 1”. The implementation passes nine regression tests including the DMG’s own worked example; reported discrepancies therefore measure the model against the rulebook’s best effort, not a strawman.

3 Methods

3.1 Features: combat math as covariates

Forty-five features in five families: raw party/roster aggregates (15); offensive potency (5); power curves and ratios (6); trait $\times$ composition interactions such as *magic resistance* $\times$ *arcane party* (6); and combat mathematics plus the DMG budget (13) — to-hit probabilities on both sides from the actual d20 rule, rounds-to-kill in both directions with a log-ratio, nova pressure (*burst_vs_pc_hp*), and the multiplier-adjusted XP budget. All feature engineering lives inside a serialized `sklearn` transformer, so training and serving cannot drift.

3.2 Model and calibration

The production model is XGBoost [3] (300 trees, depth 4, learning rate 0.149, subsample 0.68, column subsample 0.85, $\lambda=0.105$, $\alpha=0.091$, $\gamma=0.009$; tuned by Optuna [4] against a grouped-CV objective, study persisted) under **45 monotone constraints**: party level, size and role flags positive; monster count, CR, HP, DPR, burst, attack bonus, save DC and every hostile trait negative. Constraints double as the out-of-distribution guard: inputs are also clipped to legal 5e bands before featurization. Probabilities are calibrated with Platt scaling [5] on *campaign-grouped*

fold. Isotonic calibration was evaluated and rejected: its piecewise-constant map collapsed the level search (one Lich $\equiv$ two Liches) while measuring marginally worse (Brier 0.1394 vs. 0.1397).

3.3 Validation protocol

Encounters within a campaign share party, DM and house rules, so the campaign is the unit of leakage. Everything is grouped by it: `StratifiedGroupKFold` cross-validation, a grouped 80/20 holdout ($n_{\text{test}}=7,106$), and the calibration folds themselves. Headline metrics carry 95% percentile-bootstrap CIs (2,000 resamples). A kernel two-sample test (unbiased MMD² [7], RBF with median-heuristic bandwidth, permutation null) between train- and held-out-campaign feature distributions gives $\text{MMD}^2=3.1 \times 10^{-4}$, $p=0.12$: no detectable covariate shift, so the grouped-CV/holdout gap reflects campaign-level outcome variance rather than distribution drift.

4 Where the data lies: guards and simulation

4.1 DM mercy

An early model rated *nineteen* CR-21 Liches as beatable by a level-8 party. The cause is in the world, not the code: in fights the deterministic damage race declares hopeless (party deleted in ≤ 1 round), the logs still record a party win **84.5%** of the time (659 rows) — fudged rolls, retreats recorded as wins, reinforcements. The logs answer “what happens at real tables”; the application asks “what if you fought *to the death*”. That is an estimand mismatch no amount of identically-collected data can repair.

4.2 Three guards

(1) *Shape priors and clipping* (above). (2) *Roster dominance*: count-weighted averages dilute — a Lich with 2 ogres and 6 goblins has mean CR 2.94 versus the Lich’s 21, and the raw model rated the larger fight easier (level 3.25 vs. 5.0); the engine therefore also scores every homogeneous sub-roster and returns the *minimum* win probability, enforcing the axiom that adding monsters can never help the party. (3) A *survival physics cap* $P(\text{win}) \leq \sigma(a \cdot \text{TTK} + b)$, where TTK is the deterministic rounds-to-kill-the-party computed from the same feature mathematics the model trains on — smooth, monotone in level and decreasing in roster damage, so the binary search and dominance guarantees survive.

4.3 Calibrating the guard by simulation

Hand-tuned constants are a confession, so we replaced them. We vendored the public Monte Carlo engine of Battlecast (a 5e combat simulator resolving initiative, movement, spell slots and conditions with scripted




Figure 1: Survival guard vs. deathmatch truth: 180 simulated cells, the calibrated fit (production) and the previous hand-tuned curve.

play and *no mercy*) and ran a designed grid of **304 cells** $\times$ **2,000 trials** = **490,640 battles** (party fixed to Fighter/Cleric/Wizard/Rogue): a 180-cell guard grid (boss CR $\{2, 5, 10, 15, 21\} \times \text{count } \{1..19\} \times \text{level } \{1..20\}$), a 100-cell “mercy” grid (25 SRD monsters $\times$ 4 levels), and a 24-cell OOD grid (HP $\times \{1, 5, 20\}$, AC $+ \{0, 8\}$ clones). A binomially-weighted logistic fit of simulated $P(\text{win})$ on TTK gives

$$\sigma(1.63 \text{ TTK} - 3.98) \text{ replacing } \sigma(2.197 \text{ TTK} - 4.394),$$

i.e. caps of 9%/33%/71% at 1/2/3 rounds of survival (the hand guess was too generous in the 2–3-round zone), inert above ≈ 3.7 rounds where training data is trustworthy (Fig. 1). The deployed Lich ladder becomes $1 \rightarrow 3.25$, $2 \rightarrow 6.75$, $4 \rightarrow 12.75$, $8+ \rightarrow \text{beyond deadly}$ (19 Liches: $P(\text{win})$ at level 20 = 9.1%).

Pairing both lenses on the mercy grid quantifies the bias (Fig. 2): correlation 0.845, with a *crossover* — the model sits $\approx +0.32$ above the simulator on hard fights (mercy inflation; few cells, directional) and ≈ -0.14 below on easy ones (the 83% ceiling from retreats and noisy labels). On the OOD grid the extremes agree exactly (HP $\times 20$ clones win 0.000–0.003 of deathmatches) while fine mid-range ordering is only moderate (Spearman $\rho=0.54$). Caveats: the simulator implements 2024-SRD rules versus 2014-era logs, and scripted optimal play makes deathmatch rates an upper bound.

5 Results

5.1 Held-out performance

Table 1 reports performance on held-out campaigns. Calibration is tight where the data mass lives (predicted 0.8–0.9), with mild overconfidence in the lowest bins (Fig. 3). Native TreeSHAP [10] ranks party level, monster size, nova pressure (*burst_vs_pc_hp*), action economy (*num_monsters_total*) and party hit chance as the leading signals — the parsed offensive statistics the official bestiary lacks are first-class drivers.




Figure 2: DM mercy quantified: model (table reality) vs. simulator (deathmatch), one point per encounter, color = CR.

Metric (held-out, $n=7,106$)	Value	95% CI
ROC-AUC	0.657	[0.638, 0.674]
Brier score	0.139	[0.134, 0.145]
PR-AUC	0.874	—
Log loss	0.447	—
Grouped 4-fold CV AUC	0.601 ± 0.040	
Base rate $P(\text{win})$	0.830	

Table 1: Production model, campaign-grouped evaluation.

5.2 Prediction $\neq$ decision

Under identical grouped CV (Table 2), an ℓ_2 -penalized logistic regression *beats* the production model on observational risk. Promoted experimentally to production, it rated 8 Liches as trivial and a 10,000-HP monster as beatable: with no shape constraints, the linear fit absorbs selection confounding — its coefficients give *num_monsters*, *total_dpr* and *burst* *positive* signs and *avg_party_level* a *negative* one (“stronger parties lose more”, because stronger parties are served harder content). The application’s queries are interventional (“same monster, more of them”), so we deploy the constrained model and accept ≈ 0.04 AUC as the price of decision-grade behavior. The 13-axiom behavioral suite (count monotonicity, roster dominance, level monotonicity, boundary verdicts, OOD flow) now gates every model change.

5.3 Auxiliary analyses

Gaussian process. On the 797-monster CR-prediction task, a GP regressor [8] (RBF + white kernel, marginal-likelihood fitting) reaches MAE 0.86 vs. XGBoost’s 0.94 ($R^2=0.954$) with 95% predictive intervals achieving **0.95 empirical coverage**. **Synthetic balancing.** CTGAN [9]-generated synthetic




Figure 3: Reliability on held-out campaigns.

Model (identical grouped CV)	AUC $\pm$ sd	Brier
Ridge logistic	0.655 ± 0.031	0.136
RFF kernel logistic [6]	0.624 ± 0.013	0.142
XGBoost + constraints (ours)	0.614 ± 0.033	0.141

Table 2: Course model zoo. The linear model wins the metric and loses the job.

losses (leakage-guarded: fitted on training campaigns only) used to “balance” the 83/17 classes degrade the same holdout to AUC 0.638 and Brier **0.186**, dragging the predicted base rate to 0.61: with proper losses and calibration, imbalance was never the problem, and resampling destroys calibration to fix it.

6 The application

A Streamlit appraiser (auto-deployed from `main`) exposes three modes — official bestiary (762 monsters with offense provenance badges), homebrew appraiser (a CR predictor for “what the book would print”, MAE 0.9, plus the True Lethality Level), and a mixed-roster encounter builder aggregated exactly like training. Every appraisal shows the by-the-book verdict next to the model’s, win curves for five party compositions, a party-size $\times$ level heatmap, the fairest matchups with official DMG tiers, and threat flags. The target win rate is a user-facing slider: “fair” is a policy, not a constant. Engineering: 61 unit tests + 13 behavioral axioms (`make verify`), append-only experiment log, pinned environment, one-command pipeline (`make retrain`).

7 Limitations

(i) *Observational ceiling*: dice variance, table tactics and unrecorded items bound discrimination; AUC 0.657 with CIs is honest, not weak, at an 83% base rate. (ii) *Estimand mismatch* is handled by an explicit, sim-calibrated guard — a documented modeling

choice, not end-to-end learning. (iii) *Simulator caveats*: 2024-SRD rules, optimized default builds, scripted play. (iv) *Thin support on hard fights*: 96/100 mercy cells are deathmatch-easy; future work includes VTT instrumentation, 30-second post-encounter forms with explicit retreat/fudge labels (turning mercy into a censoring variable), and an all-SRD replay grid.

8 Conclusions

Five lessons travel beyond the game. Validate by the unit of leakage or every number is fiction; calibrate when the product consumes probabilities, and measure it; a better-AUC model can be unusable — constraints encode the difference between prediction and decision; know your estimand, because observational data cannot answer counterfactual questions while simulation can (and can even calibrate your priors); and the rule-book deserved a fair trial — honestly implemented, it still loses to 34,907 real fights.

References

- [1] A. Zhu et al. FIREBALL: A Dataset of Dungeons and Dragons Actual-Play with Structured Game State Information. *ACL*, 2023.
- [2] Wizards of the Coast. *Dungeon Master’s Guide*, pp. 82–83, 274. 2014.
- [3] T. Chen and C. Guestrin. XGBoost: A Scalable Tree Boosting System. *KDD*, 2016.
- [4] T. Akiba et al. Optuna: A Next-generation Hyperparameter Optimization Framework. *KDD*, 2019.
- [5] J. Platt. Probabilistic Outputs for Support Vector Machines. *Advances in Large Margin Classifiers*, 1999.
- [6] A. Rahimi and B. Recht. Random Features for Large-Scale Kernel Machines. *NeurIPS*, 2007.
- [7] A. Gretton et al. A Kernel Two-Sample Test. *JMLR*, 13, 2012.
- [8] C. E. Rasmussen and C. K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, 2006.
- [9] L. Xu et al. Modeling Tabular Data using Conditional GAN. *NeurIPS*, 2019.
- [10] S. Lundberg and S.-I. Lee. A Unified Approach to Interpreting Model Predictions. *NeurIPS*, 2017.