

SAPIENZA UNIVERSITÀ DI ROMA — STATISTICAL MACHINE LEARNING

The True Lethality Engine

Predicting the real danger of D&D 5e combat
from 34,907 recorded encounters

Presenter 1 · Presenter 2 · Presenter 3 · Presenter 4

`dnd-appraising.streamlit.app` · `github.com/danpinocontrollino/project-dnd`

Roadmap

Four parts, one question: how dangerous is this monster, *really*?

Part 1 — The problem & the data

Why the official difficulty system fails; turning 153k chat log lines into a dataset; reading damage out of statblocks.

Part 2 — Model & honest evaluation

Combat-math features; calibrated XGBoost; why we validate by *campaign*; confidence intervals.

Part 3 — Guards & simulation

Where observational data lies (DM mercy); three domain guards; calibrating one of them with 490k simulated battles.

Part 4 — Benchmarks & the product

The course model zoo; a linear model that wins AUC and still loses; the live app; limits.

Part 1

The problem & the data

Presenter: NOME 1

D&D combat in 60 seconds

- A **party** (3–6 heroes, levels 1–20) fights **monsters** run by the Dungeon Master.
- Every monster has a **Challenge Rating (CR)**: the official difficulty score, 0–30.
- The rulebook (DMG p.82) estimates encounter difficulty: sum monster XP, apply a multiplier for monster count, compare with per-level thresholds.

The official promise

A CR-5 monster $\approx$ a fair fight for four level-5 heroes.

The reality

CR is a *design intention*, never validated against play.

Our question is an estimand, not a score

Target

$\eta(x) = P(\text{party wins} \mid \text{encounter } x)$, calibrated — then invert it: the **True Lethality Level** is the party level at which a balanced party of 4 reaches a target win rate (default 65%), found by binary search on η .

- Everything downstream consumes **probabilities**, so calibration is a first-class requirement, not a nicety.
- Boundary verdicts instead of fake precision: **trivial** (wins $\geq$ target even at level 1) and **beyond deadly** (below target even at level 20).

The data: FIREBALL — real play, not simulation

- 1,471 JSONL logs of real Discord games (Avrae bot) — 153,829 turn snapshots.
- Outcome inferred from terminal HP states; unresolved fights dropped.
- **34,907 resolved encounters** across **1,462 campaigns**.
- Party win base rate: **83.0%** — remember this number.

Encounter-level row

party level/size/roles × monster
roster aggregates × outcome

`parse_fireball.py` — count-weighted means, apex maxima,
pooled sums; party levels clipped to 1–20, rosters to 30

The Attack Potency problem

The bestiary has HP and AC — but nothing about how hard a monster hits

- Solution: **parse offense from the statblock text itself** (SRD JSON): +9 to hit, 12 (2d6+5), “makes *three* attacks”, DC 14, spell lists.
- Per-action segmentation separates **sustained DPR** (weapon routine) from **burst** (breath weapon, best spell — Lich: DPR 45, burst 100 via *Power Word Kill*).
- Coverage: **321/762** monsters from real statblocks; the rest imputed from the DMG’s design table by CR.

Bugs this catches

Pooled regexes over-counted: versatile weapons (“7 . . . *or* 8 two-handed”) summed both (→ 24 monsters overstated); dragon breath leaked into multiattack math (Adult Red DPR 72.6 → 58, true $\approx$ 45–48).

Playing fair with the rulebook

If we claim to beat the DMG, we must implement the DMG correctly

- The app's "by the book" column runs the **real DMG p.82 math**: total XP $\times$ count multiplier $\rightarrow$ equivalent single-monster CR (inverse XP table).
- $6 \times \text{CR-1 monsters} = 1,200 \text{ XP}$
 $\times 2 = 2,400 \approx$ **one CR-6 monster** — not "CR 1".

Why it matters

Reported discrepancies measure our model against the rulebook's *best effort* — not against a strawman.

`official_encounter_estimate()` — 9 regression tests incl. the DMG's own worked example


Part 2

Model & honest evaluation

Presenter: NOME 2

Feature engineering: put the combat math in the features

- **To-hit geometry:** party hit chance vs monster AC; monster hit chance vs estimated party AC (both from the actual d20 rule).
- **The damage race:** rounds-to-kill on both sides; lethality_log_ratio.
- **Nova pressure:** burst_vs_pc_hp — can one action delete a hero?
- Trait $\times$ composition interactions (magic resistance $\times$ arcane party, ...).
- DMG XP budget with the count multiplier — the rulebook as a feature.



TreeSHAP: level, size, burst-vs-HP, action economy lead

The model: gradient boosting with domain priors

- **XGBoost**, 45 engineered features, tuned by Optuna (grouped-CV objective): 300 trees, depth 4, heavy subsampling — shallow and regularized, as noisy logs demand.
- **45 monotone constraints** = shape priors: higher party level can never hurt; more monster DPR can never help.
- Inputs clipped to legal 5e ranges $\Rightarrow$ graceful out-of-distribution behavior.

Serving pipeline

FeatureEngineer $\rightarrow$ Imputer $\rightarrow$ XGBoost $\rightarrow$ **Platt calibration** — one pickle, feature math travels with the model.

Why Platt, not isotonic

Isotonic's step function made 1 Lich $\equiv$ 2 Liches in the level search; sigmoid is smooth *and* measured slightly better (Brier 0.1394 vs 0.1397).

Honest evaluation: the unit of leakage is the campaign


- Encounters from one campaign share party, DM, house rules $\Rightarrow$ random splits leak.
- Everything is grouped by campaign: **StratifiedGroupKFold** CV, grouped holdout, and even the **calibration folds**.
- The gap between grouped CV (0.601) and naive splits *is* the leakage, reported rather than hidden.

Held-out campaigns ($n = 7,106$ encounters), bootstrap 95% CIs

ROC-AUC	0.657	[0.638, 0.674]
Brier score	0.139	[0.134, 0.145]
PR-AUC	0.874	—

Calibration: the product consumes probabilities

- “Which level gives 65% wins?” is meaningless if 65% doesn’t mean 65%.
- On held-out campaigns the curve hugs the diagonal where the data mass lives (0.8–0.9), with mild overconfidence in the lowest bins.
- $AUC \approx 0.66$ is **honest, not weak**: the outcome is genuinely noisy (dice, tactics, unrecorded items) and the base rate is 83%.




A model that wins AUC and still loses

The single most instructive result of the project

- Ridge logistic **beats** XGBoost on observational risk: grouped-CV AUC **0.655 vs 0.614**.
- Promoted to production, it rated **8 Liches as trivial** and a 10,000-HP monster as beatable.
- The coefficients show why:
num_monsters, total_dpr, burst all *positive*; avg_party_level *negative* — selection confounding (harder content is served to stronger parties).

Lesson: the app asks *interventional* questions. Prediction $\neq$ decision.



Part 3

Where the data lies: guards & simulation

Presenter: NOME 3

The bug that started it: 19 Liches, “level 8”

- Early model: 19 CR-21 Liches → beatable by a level-8 party at 66%.
- Root cause is in the *data*, not the code: fights the damage math calls hopeless (party deleted in ≤ 1 round) were still “won” **84.5%** of the time at real tables.
- **DM mercy**: fudged rolls, retreats recorded as wins, reinforcements. The logs answer “what happens at tables” — the app asks “what if you fought *to the death*”.

Estimand mismatch

No amount of training data fixes this: the counterfactual question is simply not in the observational distribution. This motivates *explicit domain guards*.

Three guards between the model and the user

1. **Shape priors + clipping** — monotone constraints and legal-range clipping: absurd homebrew degrades gracefully.
2. **Roster dominance** — count-weighted averages dilute (Lich + 6 goblins “looks easier” than the Lich alone); we score every homogeneous sub-roster and take the **minimum** win probability: adding monsters can never help the party.
3. **Survival physics cap** — $P(\text{win}) \leq \sigma(a \cdot \text{TTK} + b)$ where TTK = deterministic rounds-to-kill-the-party. Originally hand-tuned...

...and hand-tuned constants are a weakness — so we calibrated them externally.

Battlecast: simulated deathmatches as external ground truth

- Battlecast = Monte Carlo 5e combat simulator (initiative, movement, spell slots, conditions) — **no DM mercy**, scripted optimal play.
- Engine vendored from the site's public assets → fully local & reproducible; zero load on the service; full provenance file.
- One run = 2,000 trials $\Rightarrow$ a nearly noise-free **soft label** ($SE \leq 0.011$).

Designed grids — 304 cells,
490,640 battles

guard 180: CR {2,5,10,15,21} $\times$ count {1..19} $\times$ level {1..20}


mercy 100: 25 SRD monsters $\times$ 4 levels

OOD 24: HP $\times$ {1,5,20}, AC+{0,8} clones

make battlecast — party fixed: Fighter+Cleric+Wizard+Rogue

Result 1: the guard, calibrated


- Binomial-weighted logistic fit on the 180 guard cells:
 $\sigma(1.63 \text{ TTK} - 3.98)$ replaces $\sigma(2.197 \text{ TTK} - 4.394)$.
- Survive 1 round $\rightarrow$ cap 9%; 2 rounds $\rightarrow$ **33%** (was 50%); 3 $\rightarrow$ 71% (was 90%); inert above ≈ 3.7 — the hand guess was too generous exactly in the 2–3 round zone.
- Lich ladder now: 1 $\rightarrow$ 3.25, 2 $\rightarrow$ 6.75, 4 $\rightarrow$ 12.75, **8+ $\rightarrow$ beyond deadly.**



Result 2: DM mercy, quantified

- Same encounters through both lenses: our model (table reality) vs Battlecast (deathmatch). Correlation **0.845**.
- The gap **crosses over**: $\approx +0.32$ **above** sim on hard fights (mercy inflation), ≈ -0.14 **below** on easy ones (the 83% ceiling).
- Real curated encounters are overwhelmingly deathmatch-easy — table danger is attrition, not hopeless matchups.

Few hard-fight cells: $+0.32$ is directional. OOD: extremes agree ($HP \times 20 \rightarrow P \approx 0$), mid ranking moderate (Spearman 0.54).




Part 4

Benchmarks, the product & what we learned

Presenter: NOME 4

The course toolbox, applied where each tool fits

- **Model zoo under identical grouped CV:** ridge logistic 0.655, RFF kernel logistic 0.624, XGBoost 0.614 (AUC).
- **Kernel two-sample test (MMD):** train vs held-out campaign features — $\text{MMD}^2 = 0.00031$, permutation $p = 0.12$: *no covariate shift*; the CV-holdout gap is outcome variance, validating the split design.
- **Gaussian Process** CR predictor ($n = 797$ monsters): MAE **0.86** vs XGBoost 0.94, R^2 0.954 — and 95% predictive intervals with **0.95 empirical coverage**.



RFF = Rahimi–Recht random features: the Lec-06 kernel machine
at $n \approx 28k$

One more ablation: synthetic balancing hurts

- CTGAN-generated synthetic losses (leakage-guarded: fitted on training campaigns only) to “balance” the 83/17 classes.
- Same holdout, same pipeline: real data **AUC 0.657** / **Brier 0.139** vs GAN-balanced **0.638** / **0.186**.
- Balancing shifts the predicted base rate to 0.61 — it *destroys calibration* to fix a problem proper scoring rules never had.

Takeaway

With a proper loss and calibration, class imbalance is a non-problem; resampling “fixes” cost more than they buy.

The product: a live encounter appraiser

- **Streamlit app:** official-bestiary lookup, homebrew appraiser, mixed-monster encounter builder.
- Shows both verdicts side by side: *by-the-book* (real DMG math) vs *True Lethality Level*, with win-probability curves per party composition and a size×level heatmap.
- Target win rate is a user slider — “fair” is a policy, not a constant.

Try it

```
dnd-appraising.streamlit.app
```

Engineering rigor

61 unit tests + 13 behavioral axioms (make verify); bootstrap CIs; append-only experiment log; pinned environment; one-command pipeline (make retrain).

Limits — stated, not hidden

- **Observational ceiling:** $AUC \approx 0.66$ on a noisy, 83%-win world; the model ranks encounters well but is no oracle.
- **Estimand mismatch** handled by an explicit, now sim-calibrated guard — a modeling *choice*, documented and tested, not learned end-to-end.
- **Simulator caveats:** 2024-SRD ruleset vs 2014-era logs; scripted optimal play $\Rightarrow$ deathmatch results are an upper bound.
- Thin support where it matters most (hard fights): future work = VTT instrumentation, post-encounter forms with explicit retreat/fudge labels, an all-SRD replay grid.

Five takeaways

1. **Validate by the unit of leakage** (campaigns) — or every number is fiction.
2. **Calibrate** when the product consumes probabilities; measure it (Brier, curves), don't assume it.
3. **Prediction $\neq$ decision**: a better-AUC model can be unusable; constraints encode the difference.
4. **Know your estimand**: observational data can't answer counterfactual questions — simulation can, and can even calibrate your priors.
5. **The rulebook deserved a fair trial** — and honestly implemented, it still loses to 34,907 real fights.

Grazie!

Questions welcome — or bring us a monster to appraise.

`dnd-appraising.streamlit.app` · `github.com/danpinocontrollino/project-dnd`